

MO433 - Unsupervised Learning

Image Generation by Stable Diffusion

Alexandre Xavier Falcão

Institute of Computing - UNICAMP

afalcao@ic.unicamp.br

From Diffusion Models to Stable Diffusion

We learned diffusion models that work directly in **pixel space**.

Challenge with pixel-space diffusion

High-resolution images (e.g., $512 \times 512 \times 3$) require enormous computational resources.

- ▶ Memory: $\sim 768K$ dimensions per image.
- ▶ Training: Days/weeks on multiple high-end GPUs.
- ▶ Inference: Slow generation (1000 denoising steps).

From Diffusion Models to Stable Diffusion

We learned diffusion models that work directly in **pixel space**.

Challenge with pixel-space diffusion

High-resolution images (e.g., $512 \times 512 \times 3$) require enormous computational resources.

- ▶ Memory: $\sim 768K$ dimensions per image.
- ▶ Training: Days/weeks on multiple high-end GPUs.
- ▶ Inference: Slow generation (1000 denoising steps).

Solution: **Stable Diffusion** (Rombach et al., 2022) perform diffusion in a **compressed latent space** instead of pixel space!

Pixel Space ($512 \times 512 \times 3$) $\xrightarrow{\text{VAE}}$ Latent Space ($64 \times 64 \times 4$).

Result: $48\times$ smaller dimensionality \Rightarrow Much faster training and inference!

Stable Diffusion = Latent Diffusion + Text Conditioning

Component 1: VAE (Variational Autoencoder)

Encoder: Compress images to latent space

Image $x \in \mathbb{R}^{H \times W \times 3}$

\downarrow Conv layers (stride 2, multiple times)

Features $h \in \mathbb{R}^{h \times w \times C}$ (e.g., $h = \frac{H}{8}, w = \frac{W}{8}$)

\downarrow Flatten + Linear projections

$\mu \in \mathbb{R}^d$ (mean vector)

$\log \sigma^2 \in \mathbb{R}^d$ (log variance vector)

Reparameterization trick to sample latent:

$$z_0 = \mu + \sigma \odot \epsilon, \quad \epsilon \sim \mathcal{N}(0, I)$$

Then reshape z_0 back to spatial: $z_0 \in \mathbb{R}^{h \times w \times 4}$

Decoder: Reconstruct from latents: $\mathbb{R}^{h \times w \times 4} \rightarrow \mathbb{R}^{H \times W \times 3}$

Stable Diffusion = Latent Diffusion + Text Conditioning

Component 2: UNet (Denoising Network)

- ▶ Operates in **latent space** (not pixel space!)
- ▶ Predicts noise: $\epsilon_\theta(z_t, t, c)$
 - ▶ Input: Noisy latent z_t at timestep t
 - ▶ Output: Predicted noise ϵ
- ▶ Conditioned on text via **cross-attention**
- ▶ **Much faster** than pixel-space diffusion ($8\times$ smaller!)

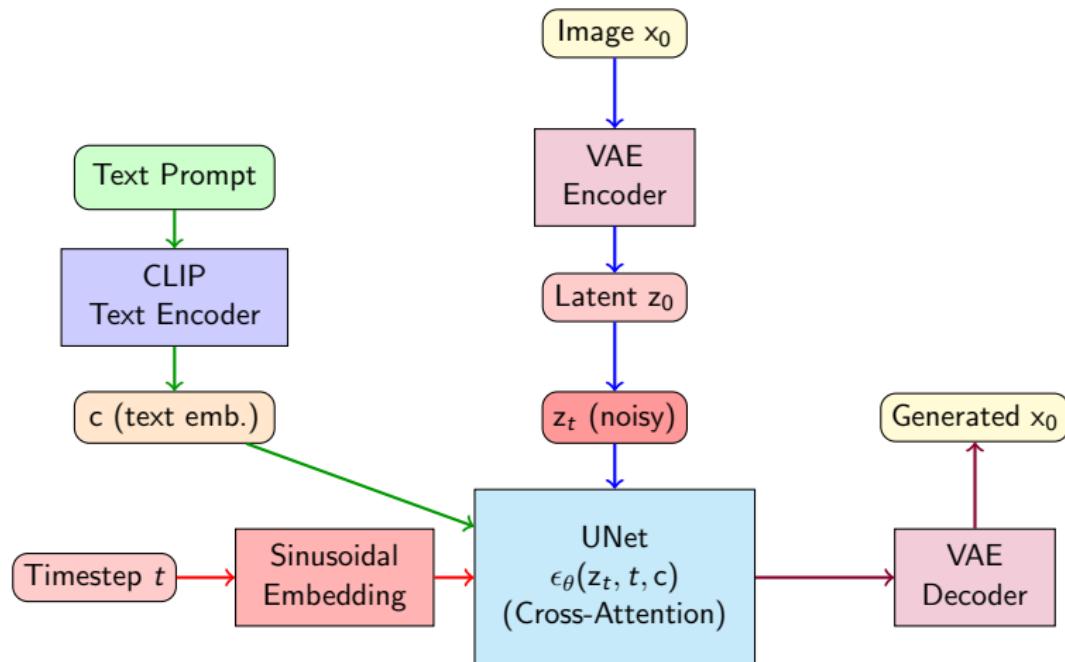
Component 3: CLIP Text Encoder

- ▶ Converts text prompts to embeddings:

$$c = \text{CLIP}(\text{"a photo of a royal guard"}) \in \mathbb{R}^{77 \times 768}$$

- ▶ Guides generation through cross-attention in UNet
- ▶ Enables text-to-image synthesis

Stable Diffusion: Architecture Overview



Training: Compress to latents, add noise, learn $\epsilon_\theta(z_t, t, c)$ to predict noise.

Inference: Start from $z_T \sim \mathcal{N}(0, I)$, iteratively denoise guided by text and timestep.

Why Stable Diffusion Works Better

Pixel-Space Diffusion:

Disadvantages:

- ▶ High dimensional (786,432D)
- ▶ Slow training
- ▶ Slow inference
- ▶ Expensive memory

Stable Diffusion:

Advantages:

- ▶ Low dimensional (16,384D)
- ▶ Fast training (48× faster)
- ▶ Fast inference
- ▶ Memory efficient
- ▶ High quality via VAE

Insight

The VAE learns to compress images into a **perceptually meaningful** latent space, discarding imperceptible details while preserving semantic information.

Agenda

- ▶ VAE: Compressing images to latent space.
- ▶ Text Conditioning with CLIP.
- ▶ Cross-Attention: Injecting text into UNet.
- ▶ Training Stable Diffusion.
- ▶ Inference and Sampling.
- ▶ Applications and fine-tuning.

VAE: Compressing Images to Latent Space

Variational Autoencoder (VAE) compresses images while preserving perceptual information.

Encoder ϕ : Maps image to latent distribution

$$q_{\phi}(z|x) = \mathcal{N}(z; \mu_{\phi}(x), \sigma_{\phi}^2(x)I)$$

Sample latent: $z = \mu_{\phi}(x) + \sigma_{\phi}(x) \odot \epsilon, \quad \epsilon \sim \mathcal{N}(0, I)$

Decoder ψ : Reconstructs image from latent

$$p_{\psi}(x|z) = \mathcal{N}(x; \mu_{\psi}(z), I)$$

Training objective (negative ELBO - Evidence Lower Bound):

$$\mathcal{L}_{\text{VAE}} = \underbrace{\|x - \hat{x}\|^2}_{\text{reconstruction}} + \underbrace{\beta \cdot \text{KL}(q_{\phi}(z|x) \| p(z))}_{\text{regularization}}$$

where $p(z) = \mathcal{N}(0, I)$ and β controls the trade-off.

VAE Architecture for Stable Diffusion

Scaling factor: Latents are scaled by $s = 0.18215$ for numerical stability.

Encoder: $\phi : \mathbb{R}^{512 \times 512 \times 3} \rightarrow \mathbb{R}^{64 \times 64 \times 4}$

- ▶ Downsampling factor: $8 \times$ in spatial dimensions.
- ▶ Compression ratio: $48 \times$ (from 786,432 to 16,384 dimensions).
- ▶ Architecture: Convolutional layers with residual blocks.

Encoding process:

$$x \in [-1, 1]^{512 \times 512 \times 3} \quad (\text{normalized image}).$$

$$x' = \phi(x) \in \mathbb{R}^{64 \times 64 \times 4}.$$

$$[\mu_\phi(x), \log \sigma_\phi^2(x)] \leftarrow [Linear(x'), Linear(x')]$$

$$z \leftarrow s \cdot [\mu_\phi(x) + \sigma_\phi(x) \odot \epsilon] \quad (\text{for diffusion}).$$

Decoder: $\psi : \mathbb{R}^{64 \times 64 \times 4} \rightarrow \mathbb{R}^{512 \times 512 \times 3}$

$$\hat{x} = \psi(z/s) \in [-1, 1]^{512 \times 512 \times 3}.$$

VAE Properties

Properties that make VAE suitable for diffusion:

1. Perceptual compression:

- ▶ Removes imperceptible high-frequency details.
- ▶ Preserves semantic and structural information.
- ▶ Low reconstruction error: $\|x - \hat{x}\|^2 < 0.01$.

2. Smooth latent space:

- ▶ Similar images map to nearby latents.
- ▶ Enables smooth interpolation.
- ▶ Suitable for Gaussian diffusion.

3. Computational efficiency:

- ▶ 48× fewer dimensions than pixels.
- ▶ Faster forward/backward passes in UNet.
- ▶ Lower memory requirements.

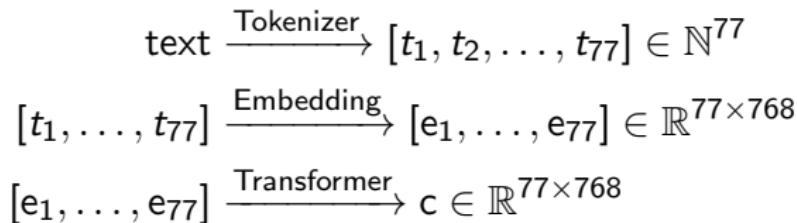
Trade-off: Slight loss of fine details (assuming large training set)
vs. massive speedup.

Text Conditioning with CLIP

CLIP (Contrastive Language-Image Pre-training) encodes text into semantic embeddings.

Architecture: Transformer-based text encoder, pre-trained on 400M image-text pairs, maps text to fixed-dimensional space: $\mathbb{R}^{77 \times 768}$.

Text encoding process:



Properties:

- ▶ Maximum sequence length: 77 tokens.
- ▶ Padding/truncation for variable-length prompts.
- ▶ Similar meanings \Rightarrow similar embeddings.

Classifier-Free Guidance

Problem: How to control the strength of text conditioning?

Solution: Classifier-Free Guidance (CFG) - interpolate between conditional and unconditional predictions.

Training: Randomly drop text conditioning ($p = 0.1$)

$$\epsilon_{\theta}(z_t, t, c) \quad \text{and} \quad \epsilon_{\theta}(z_t, t, \emptyset)$$

Inference: Combine both predictions

$$\tilde{\epsilon}_{\theta} = \underbrace{\epsilon_{\theta}(z_t, t, \emptyset)}_{\text{unconditional}} + w \cdot \underbrace{(\epsilon_{\theta}(z_t, t, c) - \epsilon_{\theta}(z_t, t, \emptyset))}_{\text{guidance direction}}$$

where $w > 0$ is the **guidance scale**.

Effects of guidance scale:

- ▶ $w = 1$: Standard conditional generation.
- ▶ $w > 1$: Stronger adherence to prompt (typical: $w = 7.5$).
- ▶ $w < 1$: Weaker conditioning, more creative.

UNet with Cross-Attention

Modified ResBlock with cross-attention:

$$h^{(0)} = \text{ResBlock}(z_t, e_t) \quad (\text{spatial conv} + \text{time embed})$$

$$h^{(1)} = \text{SelfAttn}(h^{(0)}) + h^{(0)} \quad (\text{self-attention})$$

$$h^{(2)} = \text{CrossAttn}(h^{(1)}, c) + h^{(1)} \quad (\text{text conditioning})$$

$$h^{(3)} = \text{FFN}(h^{(2)}) + h^{(2)} \quad (\text{feed-forward})$$

Three inputs to UNet:

1. z_t : Noisy latent at timestep t .
2. t : Timestep (via sinusoidal embedding e_t).
3. c : Text embeddings from CLIP.

Output: Predicted noise $\epsilon_\theta(z_t, t, c)$.

Key: Apply attention only at **lower resolutions** ($\leq 32 \times 32$ pixels).

Attention Mechanism: Core Concept

Goal: Allow each position to attend to relevant information.

Single-Head Attention:

$$Q = ZW_Q \in \mathbb{R}^{n \times d_k} \quad (\text{queries})$$

$$K = ZW_K \in \mathbb{R}^{n \times d_k} \quad (\text{keys})$$

$$V = ZW_V \in \mathbb{R}^{n \times d_v} \quad (\text{values})$$

Typically: $d_k = d_v$.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

- ▶ $QK^T \in \mathbb{R}^{n \times n}$: attention scores (who attends to whom).
- ▶ Softmax: normalize scores to probabilities.
- ▶ Multiply by V: weighted combination of values.

Limitation: Single attention pattern solved by **multiple attention heads**.

Multi-Head Attention

Run h attention operations in parallel (different “heads”).

For each head $i = 1, \dots, h$:

1. **Project to head i** using independent W_Q^i, W_K^i, W_V^i .

$$Q_i = ZW_Q^i \in \mathbb{R}^{n \times d_k}$$

$$K_i = ZW_K^i \in \mathbb{R}^{n \times d_k}$$

$$V_i = ZW_V^i \in \mathbb{R}^{n \times d_v}$$

2. **Compute attention:**

$$\text{head}_i = \text{softmax} \left(\frac{Q_i K_i^T}{\sqrt{d_k}} \right) V_i \in \mathbb{R}^{n \times d_v}$$

Combine all heads:

$$\text{MultiHead}(Z) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O \in \mathbb{R}^{n \times d}$$

where $W^O \in \mathbb{R}^{(h \cdot d_v) \times d}$ mixes information across heads - e.g., $d_k = d_v = d/h$ so concatenation gives **dimension d** (from the model).

Self-Attention in UNet: Spatial Coherence

Goal: Allow spatial locations (pixels) to attend to each other.

Self-Attention: Q, K, V all come from the same source Z.

$$Q_i = ZW_Q^i \in \mathbb{R}^{(H \cdot W) \times d_k} \quad (\text{queries from latents})$$

$$K_i = ZW_K^i \in \mathbb{R}^{(H \cdot W) \times d_k} \quad (\text{keys from latents})$$

$$V_i = ZW_V^i \in \mathbb{R}^{(H \cdot W) \times d_v} \quad (\text{values from latents})$$

where $Z \in \mathbb{R}^{(H \cdot W) \times d}$ are flattened UNet features.

Each head i :

$$\text{head}_i = \text{softmax} \left(\frac{Q_i K_i^T}{\sqrt{d_k}} \right) V_i \in \mathbb{R}^{(H \cdot W) \times d_v}$$

Output: $\text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O \in \mathbb{R}^{(H \cdot W) \times d}$

Key: $Q_i K_i^T \in \mathbb{R}^{(H \cdot W) \times (H \cdot W)}$ — each pixel attends to **all pixels!**

Cross-Attention in UNet: Text Conditioning

Goal: Condition each spatial location on the text prompt.

Cross-Attention: Q from latents, K,V from text.

$$Q_i = ZW_Q^i \in \mathbb{R}^{(H \cdot W) \times d_k} \quad (\text{queries from latents})$$

$$K_i = CW_K^i \in \mathbb{R}^{77 \times d_k} \quad (\text{keys from text})$$

$$V_i = CW_V^i \in \mathbb{R}^{77 \times d_v} \quad (\text{values from text})$$

where $Z \in \mathbb{R}^{(H \cdot W) \times d}$ (UNet), $C \in \mathbb{R}^{77 \times 768}$ (CLIP text).

Each head i :

$$\text{head}_i = \text{softmax} \left(\frac{Q_i K_i^T}{\sqrt{d_k}} \right) V_i \in \mathbb{R}^{(H \cdot W) \times d_v}$$

Output: $\text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O \in \mathbb{R}^{(H \cdot W) \times d}$

Key: $Q_i K_i^T \in \mathbb{R}^{(H \cdot W) \times 77}$ — each pixel attends to **text tokens!**

Implementation Note: Efficient Multi-Head Attention

Conceptual: h independent projections W_Q^1, \dots, W_Q^h .

Efficient implementation: Combine into single large matrix.

$$W_Q = [W_Q^1 \quad W_Q^2 \quad \cdots \quad W_Q^h] \in \mathbb{R}^{d \times (h \cdot d_k)}$$

Steps:

1. **Single projection:** $Q = ZW_Q \in \mathbb{R}^{n \times (h \cdot d_k)}$.
2. **Reshape:** Split last dimension into heads: $\mathbb{R}^{n \times h \times d_k}$.
3. **Transpose:** $\mathbb{R}^{h \times n \times d_k}$ for parallel computation.
4. **Attention:** Compute h attentions in parallel (batched).
5. **Concatenate:** Reshape back to $\mathbb{R}^{n \times (h \cdot d_k)}$.
6. **Output projection:** W^O .

Result: Mathematically equivalent, computationally efficient in GPUs!

Training Stable Diffusion

Training objective: Learn to predict noise in latent space.

Loss function:

$$\mathcal{L} = \mathbb{E}_{z_0, c, \epsilon, t} [\|\epsilon - \epsilon_\theta(z_t, t, c)\|^2]$$

where:

- ▶ $z_0 = \phi(x_0)$: Encoded latent from image.
- ▶ $c = \text{CLIP}(\text{prompt})$: Text embedding.
- ▶ $t \sim \text{Uniform}(1, T)$: Random timestep.
- ▶ $\epsilon \sim \mathcal{N}(0, I)$: Random noise.
- ▶ $z_t = \sqrt{\bar{\alpha}_t} z_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$: Noisy latent.

Frozen components:

- ▶ VAE encoder and decoder.
- ▶ CLIP text encoder.

Trainable component:

- ▶ UNet denoising network only.

Training Algorithm

Algorithm 1 Training Stable Diffusion

Require: Dataset of image-text pairs $\{(x^{(i)}, \text{prompt}^{(i)})\}$, VAE encoder ϕ , CLIP text encoder, UNet ϵ_θ .

- 1: Freeze VAE and CLIP text encoders.
- 2: **repeat**
- 3: Sample batch $\{(x^{(i)}, \text{prompt}^{(i)})\}_{i=1}^B$.
- 4: **// Encode to latent space.**
- 5: $z_0^{(i)} \leftarrow s \cdot \phi(x^{(i)})$ for all i {VAE encoding}.
- 6: **// Encode text prompts.**
- 7: $c^{(i)} \leftarrow \text{CLIP}(\text{prompt}^{(i)})$ for all i {Text encoding}.
- 8: **// Randomly drop conditioning (10% of samples).**
- 9: With probability 0.1: $c^{(i)} \leftarrow \emptyset$ {For CFG}.
- 10: **// Sample random timesteps and noise.**
- 11: Sample $t^{(i)} \sim \text{Uniform}(\{1, \dots, T\})$ for all i .
- 12: Sample $\epsilon^{(i)} \sim \mathcal{N}(0, I)$ for all i .
- 13: **// Add noise to latents.**
- 14: $z_t^{(i)} \leftarrow \sqrt{\bar{\alpha}_{t^{(i)}}} z_0^{(i)} + \sqrt{1 - \bar{\alpha}_{t^{(i)}}} \epsilon^{(i)}$.
- 15: **// Predict noise.**
- 16: $\hat{\epsilon}^{(i)} \leftarrow \epsilon_\theta(z_t^{(i)}, t^{(i)}, c^{(i)})$ {UNet forward}.
- 17: **// Compute loss and update.**
- 18: $\mathcal{L} \leftarrow \frac{1}{B} \sum_{i=1}^B \|\epsilon^{(i)} - \hat{\epsilon}^{(i)}\|^2$.
- 19: Update θ using $\nabla_\theta \mathcal{L}$.
- 20: **until** converged.

Sampling Algorithm (Inference)

Algorithm 2 Sampling from Stable Diffusion

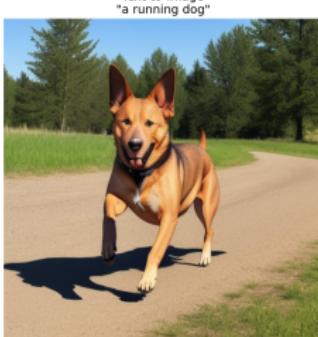
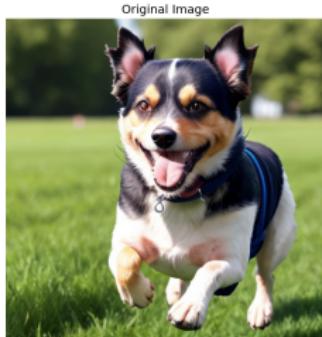
Require: Trained UNet ϵ_θ , VAE decoder ψ , CLIP text encoder, text prompt, guidance scale w

```
1: // Encode text prompt
2:  $c \leftarrow \text{CLIP}(\text{prompt})$  {Conditional embedding}
3:  $c_\emptyset \leftarrow \text{CLIP}("")$  {Unconditional (empty) embedding}
4: // Start from random noise
5: Sample  $z_T \sim \mathcal{N}(0, I)$  {Pure noise in latent space}
6: // Denoising loop
7: for  $t = T, T - 1, \dots, 1$  do
8:   // Predict noise with and without conditioning
9:    $\epsilon_{\text{uncond}} \leftarrow \epsilon_\theta(z_t, t, c_\emptyset)$ 
10:   $\epsilon_{\text{cond}} \leftarrow \epsilon_\theta(z_t, t, c)$ 
11:  // Apply classifier-free guidance
12:   $\hat{\epsilon} \leftarrow \epsilon_{\text{uncond}} + w \cdot (\epsilon_{\text{cond}} - \epsilon_{\text{uncond}})$ 
13:  // Denoise one step
14:   $z_{t-1} \leftarrow \text{scheduler.step}(\hat{\epsilon}, t, z_t)$ 
15: end for
16: // Decode to image
17:  $x_0 \leftarrow \psi(z_0 / s)$  {VAE decoding}
18: return  $x_0$ 
```

Inference: Visual Demonstration

Text-to-Image Generation (see `code1-stable-diffusion.py`)

Prompt: “*a running dog*”



Left: Original image for reference.

Center: Text-to-image (from pure noise).

Right: Image-to-image (original + noise, then denoise with text guidance).

Image-to-Image Generation

Start from a noisy version of an existing image to preserve its structure rather than generating a new image from pure noise.

Process:

1. Encode source image: $z_0 = \phi(x_{\text{source}})$.
2. Add controlled noise:

$$z_{t_0} = \sqrt{\bar{\alpha}_{t_0}} z_0 + \sqrt{1 - \bar{\alpha}_{t_0}} \epsilon$$

3. Denoise from t_0 to 0 (instead of T to 0).
4. Decode: $x_{\text{output}} = \psi(z_0)$.

Strength parameter $s \in [0, 1]$, $t_0 = s \cdot T$.

- ▶ $s = 0$: No change (skip denoising).
- ▶ $s = 0.3$: Subtle modifications, preserves image structure.
- ▶ $s = 0.7$: Major changes, follows prompt more.
- ▶ $s = 1.0$: Equivalent to text-to-image.

Applications of Stable Diffusion

1. **Text-to-Image:** Generate images from text descriptions (e.g., art generation).
2. **Image-to-Image:** Transform existing images (e.g., style transfer).
3. **Inpainting:** Fill masked regions (e.g., object removal, image completion).
4. **Super-resolution:** Upscale low-res images (e.g., detail enhancement, quality improvement).

Fine-tuning Stable Diffusion

Why fine-tune?

- ▶ Adapt to specific domains (e.g., medical images, art styles).
- ▶ Learn new concepts or objects.
- ▶ Improve quality on specific types of prompts.

Full Fine-tuning:

- ▶ Train entire UNet on custom dataset.
- ▶ Requires: Many images (1000+), high compute (GPU days).
- ▶ Best quality but expensive.

Efficient Alternative: LoRA (Low-Rank Adaptation).

- ▶ Freeze original weights, add small trainable layers.
- ▶ Requires: Few images (10-100), low compute (GPU hours).
- ▶ Good quality with minimal resources.

Let's understand LoRA...

LoRA: Low-Rank Adaptation

Instead of updating all weights $W \in \mathbb{R}^{d \times k}$, add a low-rank perturbation.

Original layer:

$$y = Wx$$

LoRA layer:

$$y = \underbrace{W}_{\text{frozen}} x + \underbrace{BA}_{\text{trainable}} x$$

where:

- ▶ $W \in \mathbb{R}^{d \times k}$: Original frozen weights.
- ▶ $B \in \mathbb{R}^{d \times r}, A \in \mathbb{R}^{r \times k}$: Trainable low-rank matrices.
- ▶ $r \ll \min(d, k)$: Rank (typically $r = 4, 8, 16$).

Parameter reduction:

$$\text{Original: } d \times k \quad \Rightarrow \quad \text{LoRA: } r(d + k)$$

Ex: $d = k = 1024, r = 8: 1,048,576 \Rightarrow 16,384$ ($64 \times$ reduction!).

LoRA: Low-Rank Adaptation

Intrinsic Dimensionality Hypothesis:

It works when the updates to weights during fine-tuning lies in a low-dimensional subspace.

Mathematical intuition:

Full update: $W_{\text{new}} = W + \Delta W$

LoRA approximation: $\Delta W \approx BA$ where $\text{rank}(BA) = r$

Benefits for Stable Diffusion fine-tuning:

1. **Memory efficient:** Only train $\sim 1\%$ of parameters.
2. **Fast training:** Fewer parameters to update.
3. **Modular:** Can swap LoRA weights without changing base model.
4. **Composable:** Can combine multiple LORAs.
5. **No catastrophic forgetting:** Base model stays intact.

Where to apply LoRA: Attention projection layers in UNet.

Fine-tuning: Training Setup

Full fine-tuning (see `code2-finetune-stable-model.py`):

Frozen components:

- ▶ VAE encoder/decoder.
- ▶ CLIP text encoder.

Trainable:

- ▶ Entire UNet (or with LoRA: only LoRA layers).

Training hyperparameters:

- ▶ Learning rate: 1×10^{-5} (lower than pretraining).
- ▶ Batch size: 1-4 (limited by GPU memory).
- ▶ Resolution: 384×384 or 512×512 .
- ▶ Epochs: 10-100 (depends on dataset size).

Memory optimizations:

- ▶ Gradient checkpointing (trades compute for memory).
- ▶ 8-bit optimizer (50% memory reduction).
- ▶ Mixed precision training (FP16).

Fine-tuning Algorithm

Algorithm 3 Fine-tuning Stable Diffusion (Full or LoRA)

Require: Custom dataset $\{(x^{(i)}, \text{prompt}^{(i)})\}$, pretrained model

```
1: Load pretrained VAE, CLIP, UNet
2: Freeze VAE and CLIP
3: if using LoRA then
4:   Freeze UNet base weights
5:   Add LoRA layers to attention projections
6:   Initialize A ~  $\mathcal{N}(0, \sigma^2)$ , B = 0
7: end if
8: repeat
9:   Sample batch  $\{(x^{(i)}, \text{prompt}^{(i)})\}$ 
10:  Encode:  $z_0^{(i)} \leftarrow \mathcal{E}(x^{(i)})$ ,  $c^{(i)} \leftarrow \text{CLIP}(\text{prompt}^{(i)})$ 
11:  Sample  $t^{(i)}$ ,  $\epsilon^{(i)}$ 
12:  Compute  $z_t^{(i)} = \sqrt{\bar{\alpha}_{t(i)}} z_0^{(i)} + \sqrt{1 - \bar{\alpha}_{t(i)}} \epsilon^{(i)}$ 
13:  Predict:  $\hat{\epsilon}^{(i)} \leftarrow \epsilon_\theta(z_t^{(i)}, t^{(i)}, c^{(i)})$ 
14:  Loss:  $\mathcal{L} = \frac{1}{B} \sum_i \|\epsilon^{(i)} - \hat{\epsilon}^{(i)}\|^2$ 
15:  if using LoRA then
16:    Update only LoRA weights (A, B)
17:  else
18:    Update all UNet weights
19:  end if
20: until converged
21: Save fine-tuned model (or LoRA weights)
```

Fine-tuned Model Inference

Inference with fine-tuned model.

Process:

1. Load pretrained VAE and CLIP (unchanged).
2. Load fine-tuned UNet (or base UNet + LoRA weights).
3. Generate images using standard sampling algorithm.

Example prompt: “*a jumping dog with sunglasses wearing a bear hat*”



(Image generated using `code3-stable-model-inference.py` with user prompt)

What We Learned

- 1. Stable Diffusion = Latent Diffusion + Text Conditioning**
 - ▶ VAE compresses to $48\times$ smaller latent space.
 - ▶ Diffusion operates in latent space (much faster).
- 2. Three key components:**
 - ▶ VAE: Compression and reconstruction.
 - ▶ CLIP: Text understanding.
 - ▶ UNet: Denoising with cross-attention.
- 3. Training:** Predict noise in latent space conditioned on text.
- 4. Inference:** Iterative denoising with classifier-free guidance.
- 5. Fine-tuning:** Adapt to custom domains.
 - ▶ Full: High quality, high cost.
 - ▶ LoRA: Good quality, low cost (recommended).

See also other techniques: DreamBooth, IP-Adapter, Textual Inversion, and ControlNet.

Take-home Message

- ▶ VAE (see `code4-train-vae.py`) requires several thousands of samples at least to reduce blurring and generate clear images in stable diffusion.
- ▶ Training a simple diffusion model from scratch also requires many samples (`code5-train-diffusion-from-scratch.py`), but stable diffusion can considerably amend the problem (`code5-train-stable-diffusion-from-scratch.py`).
- ▶ Real life problems may present a few images from domain difficult to adapt the model via fine-tuning.
- ▶ LoRA is an alternative (`code6-train-model-with-lora.py` and `code7-generate-image-with-lora.py`), but

can it adapt a pretrained stable diffusion model to a distinct domain with a few images?